

CSE6140 Fall 2017 Project - Minimum Vertex Cover

November 2, 2017

1 Overview

The Minimum Vertex cover (MVC) problem is a well known NP-complete problem with numerous applications in computational biology, operations research, the routing and management of resources. In this project, we will treat the problem using different algorithms, evaluating their theoretical and experimental complexities on real world datasets.

2 Objectives

- Get hands-on experience in solving an intractable problem that is of practical importance
- Implement an exact branch-and-bound algorithm
- Implement approximate algorithms that run in reasonable time and provide high-quality solutions with guarantees
- Implement heuristic algorithms (without approximation guarantees)
- Develop your ability to conduct empirical analysis of algorithm performance on datasets, and to understand the trade-offs between accuracy, speed, etc., across different algorithms
- Develop teamwork skills while working with other students


3 Groups

Your groups have been posted to T-Square under “Resources → Project → Project Groups.pdf”.

4 Background

Given an undirected graph $G = (V, E)$ with a set of vertices V and a set of edges E , a vertex cover is a subset $C \subseteq V$ such that $\forall (u, v) \in E : u \in C \vee v \in C$. The Minimum Vertex Cover problem is therefore simply the problem of finding minimizing $|C|$.

For more details on the MVC problem, refer to Karp’s original 21 NP-Complete problems <http://www.cs.berkeley.edu/~luca/cs172/karp.pdf>. The following diagram shows an example of minimal vertex cover in light shade:



5 Algorithms

You will implement algorithms that fall into three categories:

1. exact, computing an optimal solution to the problem
2. construction heuristics some of which have approximation guarantees
3. local search with no guarantees but usually much closer to optimal than construction heuristics.

In what follows, we present the high-level idea behind the algorithms you will implement

1. **Exact algorithm using Branch-and-Bound.** Implement the Branch-and-Bound algorithm as seen in class. You may design any lower bound of your choice, but consider the following strategies as a guidance:
 - Let C' be a partial vertex cover of G , and let $G' = (V', E')$ be the subgraph of G not covered by C' . $V' = V - C' - \{v \in V | \forall (v, u) \in E, u \in C'\}$ (note we prune nodes that cannot help us cover any uncovered edges) and $E' = \{(u, v) \in E | v, u \in V'\}$. Then, a lower bound can be calculated as $LB = |C'| + (\text{lower bound on the VC for } G')$. [Note that one lower bound we know how to obtain is from our approximation algorithms (the approx solution quality A is $\leq 2OPT$, hence $LB = \frac{1}{2}A \leq OPT$). The approx algorithm in class finds a maximal matching given a graph, one way to find an even tighter bound is to compute the maximum matching of G' (this a more complicated problem but still has a polynomial time algorithm). Alternatively, for those of you who want to try this, given G' you can construct the corresponding LP relaxation for finding minimum VC, use an available LP solver in your programming language of choice, and obtain the LP optimal solution as a lower bound to use in your branch and bound implementation (note you are not allowed to use a solver that directly solved the ILP using branch-and-bound).]

Since this algorithm is still of worst-case exponential complexity, your algorithm must have additional code that allows it to stop after running for T minutes, and to return the current best solution that has been found so far. Clearly, for small datasets, this algorithm will most likely return an optimal solution, whereas it will fail to do so for larger datasets.

2. **Construction Heuristics with approximation guarantees.** Please implement a constructive heuristic for MVC with approximation guarantees as shown in class. The algorithm in class chooses any remaining edge in E' . You may decide to augment the algorithm with some ways of prioritizing the vertices in E' . Optionally, you may implement one of the approximation algorithms with guarantees described in [4].
3. **Local Search.** There are many variants of local search (LS) and you are free to select which one you want to implement. Please implement 2 types/variants of local search. They can be in different families of LS such as Simulated Annealing vs Genetic Algorithms vs Hill Climbing, or they can be in the same general family but should differ by the neighborhood they are using, or by the perturbation strategy, etc. They need to be different enough to observe qualitative differences in behavior. Optionally, you may use one of the algorithms described in [4] to compute initial solutions for your variants of local search.

Tips

While there are many approaches([1, 2, 5, 3]), a common local search framework stems from considering the problem in its decisional version: Given an integer k , is there a cover of size at most k ? To solve the k -vertex cover problem a candidate solution of size k is maintained and two vertices (2-neighborhood) or more are iteratively exchanged in and out until it becomes a vertex cover. Once a k -vertex cover is achieved, l vertices are removed and a $k - l$ vertex cover is searched instead. See Algorithm 1 for a pseudocode description of the framework. Within the framework, there are multiple parameters and strategies to be defined. These include the selection scheme for entering and exiting

vertices, as well as the pricing function that determines the quality of a given candidate solution. An example of pricing function for a given candidate solution C' is as follows:

$$Cost(G, C') = \text{number of edges not covered by } C'$$

Given a candidate solution C' , an example of selection scheme is given as follows:

- (a) Exiting vertex: Select the vertex $v \in C'$ that produces the minimum increase in cost.
- (b) Entering vertex: Select an endpoint from a random uncovered edge.

Algorithm 1 Local Search Framework

```

1: procedure LOCALSEARCH( $V, E$ )
2:    $C \leftarrow$  initial vertex cover
3:   while elapsed time < cutoff do
4:     if  $C$  is a vertex cover then
5:        $C^* \leftarrow C$ 
6:       Remove one or more than one vertices from  $C$ 
7:     end if
8:      $u \leftarrow$  Select exiting vertex from  $C$ 
9:      $C \leftarrow C \setminus \{u\}$ 
10:     $v \leftarrow$  Select entering vertex from  $V \setminus C$ 
11:     $C \leftarrow C \cup \{v\}$ 
12:  end while
13:  return  $C^*$ 
14: end procedure

```

We can also look at the vertex cover from the perspective of its dual problem: the Maximum Independent Set (MIS). An independent set of a graph $G = (V, E)$ is a subset of V whose elements are pairwise non-adjacent. Notice that finding the Minimum Vertex Cover C of a graph is equivalent to finding the Maximum Independent set M . As a matter of fact, $C = V \setminus M$. Consider the MIS problem for more local search neighborhood definitions. Here are two examples:

- (a) 2–Improvement: A vertex of a maximal independent set is replaced for two of its neighbors that are not adjacent to each other. This way, the total number of vertices in the solution is increased by one.
- (b) 3–Improvement: Given a 2–maximal solution, look for two vertices $\{x, y\} \in M$ and three vertices $\{u, v, w\} \in V \setminus M$ such that we can improve M by removing $\{x, y\}$ and adding $\{u, v, w\}$.

6 Data

You will run the algorithms you implement on some real and random datasets extracted from the 10th DIMACS challenge, which can be downloaded from T-Square: Resources \rightarrow Project \rightarrow DATA.zip

Each graph in the data set is represented in the popular Metis I/O format. An unweighted graph $G = (V, E)$ is represented as follows:

1. A first line with three integer numbers representing the size of the graph and a zero(indicating the graph is unweighted): $|V|, |E|, 0$
2. $|V|$ lines follow. Where line i specifies the list the adjacencies for vertex i .

See section 4 of <http://www.cc.gatech.edu/dimacs10/data/manual.ps> for further details. Table 1 describing each of the real-world instances.

Problem name	Description	$ V $	$ E $	Opt
jazz.graph	Jazz musicians network. List of edges of the network of Jazz musicians.	198	2742	158
karate.graph	Zachary's karate club: social network of friendships between 34 members of a karate club at a US university in the 1970s.	34	78	14
football.graph	American College football: network of American football games between Division IA colleges during regular season	115	613	94
as-22july06.graph	Graph of the whole Internet: a symmetrized snapshot of the structure of the Internet at the level of autonomous systems, reconstructed from BGP tables	22963	48436	3303
hep-th.graph	Weighted network of coauthorships between scientists	8361	15751	3926
star.graph	Star-like structures of different graphs with different types	11023	62184	6902
star2.graph	Star-like structures of different graphs with different types	14109	98224	4542
netscience.graph	Coauthorship network of scientists working on network theory	1589	2742	899
email.graph	List of edges of the network of e-mail interchanges between members of a university	1133	5451	594
delaunay n10.graph	Delaunay triangulations of random points in the unit square	1024	3056	703
power.graph	Network representing the topology of the Western States Power Grid	4941	6594	2203

Table 1: Real world instances

7 Code

All your code files should include a top comment that explains what the given file does. Your algorithms should be well-commented and self-explanatory. Use a README file to explain the overall structure of your code, as well as how to run your executable.

Your executable should take as input i) the filename of a dataset and ii) the cut-off time (in seconds) iii) the method to use, iv) a random seed. If it is run with the same 4 input parameters, your code should produce the same output. The executable must conform with the following format:

```
exec -inst <filename> -alg [BnB|Approx|LS1|LS2] -time <cutoff in seconds> -seed <random seed>
```

Any run of your executable with any four inputs (filename, cut-off time, method and seed) must produce two types of output files:

1. Solution files:

- File name: $\langle instance \rangle - \langle method \rangle - \langle cutoff \rangle - \langle randSeed \rangle *.sol$, e.g. *Atlanta_BnB_600.sol*, *jazz_LS1_600_4.sol* Note that as in the example above, randSeed is only applicable when the method of choice is randomized (e.g. local search). When the method is deterministic (e.g. branch-and-bound), randSeed is omitted from the solution file's name.
- File format:
 - (a) line 1: quality of best solution found (integer)

- (b) lines 2: list of vertex IDs of the vertex cover (comma-separated): v_1, v_2, \dots, v_n
2. Solution trace files:
- File name: $\langle \text{instance} \rangle _ \langle \text{method} \rangle _ \langle \text{cutoff} \rangle _ \langle \text{randSeed} \rangle _ \text{*.trace}$, e.g. *jazz_BnB_600.trace*, *jazz_LS1_600_4.trace*. Note that randSeed is used as in the solution files.
 - File format: each line has two values (comma-separated):
 - (a) A timestamp in seconds (double)
 - (b) Quality of the best found solution at that point in time (integer). Note that to produce these lines, you should record every time a new improved solution is found.
- Example:
- 3.45, 102
7.94, 95

8 Output

You should run all the algorithms you have implemented on all the instances we provide, and submit the output files generated by your executable, as explained in the Code section.

P.S.: DO NOT SUBMIT THE INPUT DATA FILES.

9 Evaluation

We now describe how you will use the outputs produced by your code in order to evaluate the performance of the algorithms.

1. Comprehensive Table: Include a table with columns for each of your MVC algorithms as seen below. For all algorithms report the relative error with respect to the optimum solution quality provided (soon) to you in the MVC instance files. Relative error (*RelErr*) is computed as $(Alg - OPT)/OPT$. Round time and *RelErr* to two significant digits beyond the decimal. For local search algorithms, your results for each cell should be the average of some number (at least 10) of runs with different random seeds for that dataset. You will fill in average time (seconds) and average vertex cover size. You can also report in any other information you feel is interesting.

	Branch and Bound			Etc. (other algorithms)		
Dataset	Time (s)	Length	RelErr	Time (s)	Length	RelErr
instance	3.26	3400	0.0021			

The next three evaluation plots are applicable to local search algorithms only, and need only be presented for the instances *XX* and *YY*. All the information you need to produce these plots is in your solution trace files.

1. Qualified Runtime for various solution qualities (QRTDs): A plot similar to those in Lecture on Empirical Evaluation. The x-axis is the run-time in seconds, and the y-axis is the fraction of your algorithm runs that have 'solved' the problem. Note that 'solve' here is w.r.t. to some relative solution quality q^* . For instance, for $q^* = 0.8\%$, a point on this plot with x value 5 seconds and y value 0.6 means that in 60% of your runs of this algorithm, you were able to obtain a VC of size at most the optimal size plus 0.8% of that. When you vary q^* for a few values, you obtain the points similar to those in slide 24.
2. Solution Quality Distributions for various run-times (SQDs): Instead of fixing the relative solution quality and varying the time, you now fix the time and vary the solution quality. The details are analogous to those of QRTDs.

3. Box plots for running times: Since your local search algorithms are randomized, there will be some variation in their running times. You will use box plots, as described in the ‘Theory’ section of this blog post: <http://informationandvisualization.de/blog/box-plot>. Read the blog post carefully and understand the purpose of this type of plots. You can use online box plot generators such as <http://boxplot.tyerslab.com/> to produce the plot automatically from your data.

10 Report

1. Formatting

You will use the format of the Association for Computing Machinery (ACM) Proceedings to write your report. See <http://www.acm.org/publications/proceedings-template> for LaTeX and Word downloads.

2. Content

Your report should be written as if it were a research paper in submission to a conference or journal. A sample report outline looks like this:

- Introduction: a short summary of the problem, the approach and the results you have obtained.
- Problem definition: a formal definition of the problem.
- Related work: a short survey of existing work on the same problem, and important results in theory and practice.
- Algorithms: a detailed description of each algorithm you have implemented, with pseudo-code, approximation guarantee (if any), time and space complexities, etc. What are the potential strengths and weaknesses of each type of approach? Did you use any kind of automated tuning or configuration for your local search? Why and how you chose your local search approaches and their components? Please cite any sources of information that you used to inform your algorithm design.
- Empirical evaluation: a detailed description of your platform (CPU, RAM, language, compiler, etc.), experimental procedure, evaluation criteria and obtained results (plots, tables, etc.). What is the lower bound on the optimal solution quality that you can drive from the results of your approximation algorithm and how far is it from the true optimum? How about from your branch-and-bound?
- Discussion: a comparative analysis of how different algorithms perform with respect to your evaluation criteria, or expected time complexity, etc.
- Conclusion

11 Deliverables

Failure to abide by the file naming and folder structure as detailed here will result in penalties.

1. Each student should submit the same PDF report, titled

group_ < group_num > - < gtusername > -Report.pdf

following the guidelines in section Report.

2. Each student should submit the same ZIP file, titled:

group_ < group_num > - < gtusername > .zip

Example: *group_21 - smith3.zip*

The file must have the following files/folders:

- Code: a folder named ‘code’ that contains all your code, the executable and a *README-group_<group_num>.txt* file, as explained in section Code.
 - Output: a folder named ‘output’ that contains all output files, as explained in section Code.
3. Each student should submit a *private* evaluation of their team members. For each team member (including self) include a score from 0 to 10, outline the contributions that the member did to the project, and justification for the score.

References

- [1] Diogo V. Andrade, Mauricio G.C. Resende, and Renato F. Werneck. Fast local search for the maximum independent set problem. *Journal of Heuristics*, 18(4):525–547, 2012.
- [2] Shaowei Cai. Balance between complexity and quality: local search for minimum vertex cover in massive graphs. In *Proceedings of the Twenty-Fourth International Joint Conference on Artificial Intelligence, IJCAI*, pages 25–31, 2015.
- [3] Shaowei Cai, Kaile Su, and Abdul Sattar. Local search with edge weighting and configuration checking heuristics for minimum vertex cover. *Artificial Intelligence*, 175(9):1672–1696, 2011.
- [4] François Delbot and Christian Laforest. Analytical and experimental comparison of six algorithms for the vertex cover problem. *Journal of Experimental Algorithmics (JEA)*, 15:1–4, 2010.
- [5] Shaowei Cai, Kaile Su and Abdul Sattar. Two new local search strategies for minimum vertex cover. 2012.